High-Performance BFS on Billion-Scale Graphs

V3.3: 831ms on Friendster Using CUDA

Evangelos Moschou
AEM: 10986

January 2026

Abstract

This report presents the design and optimization of a GPU-accelerated Breadth-First Search (BFS) solver capable of traversing the Friendster social network (65.6M nodes, 3.6B edges) in **831 milliseconds**—a **26x speedup** over sequential CPU. For graphs that fit entirely in VRAM (20M nodes), we achieve an even higher **27.6x speedup** (190ms). We implement the three required versions: (1) **V1: Dynamic Thread Assignment**, (2) **V2: Chunk-Based Processing**, and (3) **V3: Shared-Memory Cooperative**. We further optimize V3 into **V3.3 (Final)**, introducing: (A) **Direct Queue Emission** to eliminate O(N) scans; (B) **Hybrid Direction-Optimization**; (C) **Delta-Varint Compression**; and (D) **Zero-Copy Streaming**.

1 Introduction

Graph algorithms are fundamental to social network analysis, recommendation systems, and scientific computing. However, billion-scale graphs present significant challenges for GPU acceleration due to:

- **Memory Constraints**: Large graphs exceed GPU VRAM (Friendster: 14.4GB vs 5.8GB available).
- **Irregular Access Patterns**: Power-law degree distributions cause load imbalance.
- **Atomic Contention**: Concurrent frontier updates create serialization bottlenecks.

This project addresses these challenges through a systematic optimization of the BFS algorithm, culminating in **Version 3.3** which achieves sub-second traversal on Friendster.

1.1 Contributions

1. **Direct Queue Emission**: A novel technique where Bottom-Up kernels build the next frontier queue directly during traversal, eliminating expensive O(N) scans.
2. **Warp-Aggregated Atomics**: Reducing global atomic contention by 32x through hierarchical aggregation.
3. **Zero-Copy Streaming**: Enabling processing of out-of-core graphs via pinned host memory.
4. **Comprehensive Evaluation**: Detailed performance analysis across algorithm versions.

2 System Design

2.1 Data Representation: Delta-Compressed CSR

Standard CSR (Compressed Sparse Row) format stores:

- `row_ptr[N+1]`: Offsets into column array (8 bytes each for 64-bit indexing)
- `col_idx[E]`: Neighbor IDs (4 bytes each)

For Friendster, this requires 14.4GB—exceeding our 5.8GB VRAM. We implement **Delta-Varint Compression**:

1. Sort neighbors for each row
2. Store differences: $\Delta_i = \text{neighbor}_i - \text{neighbor}_{i-1}$
3. Encode deltas using Variable-Length Quantity (1-5 bytes per delta)

This reduces Friendster to **9.13 GB** (37% reduction), enabling Zero-Copy streaming from host RAM.

2.2 Zero-Copy Memory Management

Since compressed data still exceeds VRAM, we use `cudaHostRegister` to pin host memory and map it to device address space. Kernels access graph data over PCIe with effective bandwidth of ∼1.4 GB/s. The key insight is that BFS traversal exhibits high temporal locality per level—we stream each portion of the graph once per BFS iteration.

2.3 Hybrid Direction-Optimization

We implement the Beamer direction-optimization strategy:
- **Top-Down**: Process small frontiers by expanding each frontier node's neighbors.
- **Bottom-Up**: Process large frontiers by checking if each unvisited node has a frontier neighbor.

The switching threshold is $N/26$ (≈2.5M nodes for Friendster). This hybrid approach reduces redundant edge checks by up to 50% on high-diameter graphs.

3 V3.3 Optimizations (Advanced V3)

3.1 Direct Queue Emission (Critical)

In prior versions, Bottom-Up traversal marked newly discovered nodes by writing distances, then required a separate O(N) scan (`distancesToQueueKernel`) to rebuild the frontier queue. This scan dominated runtime on large graphs.

V3.3 Solution: Bottom-Up kernels now emit discovered nodes directly to the next frontier using warp-aggregated atomics:
1. Each warp processes one unvisited node
2. If a frontier neighbor is found, thread sets `found = true`
3. Warp ballot collects all `found` flags: `__ballot_sync()`
4. Leader thread atomically reserves space: `atomicAdd(queue_size, popcount)`
5. Each found thread writes its node ID at the computed offset

This eliminates the O(N) scan entirely, reducing Bottom-Up level time by 100–150ms.

3.2 Warp-Aggregated Atomics (Top-Down)

Top-Down kernels originally used per-discovery `atomicAdd` calls to append to the next frontier. With billions of edges, this created severe contention.

V3.3 Solution: Within each warp, discoveries are aggregated before a single atomic operation:
1. Each thread processes edges and flags discoveries
2. `__ballot_sync()` + `__popc()` count discoveries per warp
3. Leader performs one `atomicAdd` for the entire warp
4. `__shfl_sync()` broadcasts base index to all threads
5. Each thread computes its local offset and writes

This reduces global atomics from O(edges) to O(warps), a 32x reduction.

3.3 Synchronization Reduction

We removed unnecessary `cudaDeviceSynchronize()` calls after kernel launches, relying instead on CUDA's implicit stream serialization. This reduced per-level overhead by ∼10%.

4 Experimental Results

4.1 Test Environment

- **GPU**: NVIDIA GeForce RTX 3060 Laptop (Ampere, sm_86, 5.8GB VRAM)
- **System**: 32GB RAM, CUDA 12.5, Ubuntu Linux
- **Compiler**: nvcc -O3 -arch=sm_86

4.2 Benchmark: Friendster Social Network

- Vertices: 65,608,366
- Edges: 3,612,134,270 (3.6 billion)
- Uncompressed size: 14.4 GB
- Compressed size: 9.13 GB (1.58x ratio)
- Diameter: 22 levels

4.3 Performance Evolution

Version	Key Optimization	Time (ms)	Speedup
Sequential CPU	STL Queue + Vector	21,708	1.0x
V1 Dynamic	Atomic Work Queue	21,245	1.02x
V2 Chunked	Thread-Per-Chunk	18,506	1.2x
V3.1 Shared	Shared Memory Kernels	12,009	1.8x
V3.2 Hybrid	Direction-Optimization	1,200	18.1x
V3.3 Final	**Direct Queue + Compression**	**831**	**26.1x**

Table 1: BFS Performance Evolution on Friendster (source=0)

Metric	BFS V3.3	Afforest
Runtime	**831 ms**	21.6 s
Edges Processed	3.6 B	3.6 B
Throughput (GTEPS)	4.35	0.17
Memory Streamed	9.13 GB	9.13 GB
Effective Bandwidth	11.0 GB/s	0.42 GB/s

Table 2: Detailed metrics. GTEPS = Giga Traversed Edges Per Second.

4.4 Analysis

The 26.1x speedup over CPU is achieved through:

- **V3.1 Shared**: 1.8x speedup using shared memory, but limited by global atomics.
- **V3.2 Hybrid**: 10x jump by skipping redundant checks (Direction-Optimization).
- **V3.3 Final**: Further 30% improvement via Compression and Direct Queue Emission.

It is important to note that the absolute performance (and thus the speedup) on Friendster is strictly bounded by the PCIe bus bandwidth. The compressed graph occupies **9.13 GB**, which exceeds the GPU's **5.8 GB VRAM**. Consequently, the GPU must stream data from system RAM during traversal.

4.5 VRAM-Resident Graph Analysis

To quantify the algorithmic speedup without the PCIe bottleneck, we benchmarked a synthetic random graph (20M nodes, 200M edges, 2.0GB total) that fits entirely in VRAM.

Version	Time (ms)	Speedup	Observation
Sequential CPU	5,278	1.0x	Baseline
V2 Chunked	475	11.1x	High serialization overhead
V3.3 Compressed	381	13.8x	Decompression overhead dominates
V1 Dynamic	279	18.9x	Low overhead, massive parallelism
V3.1 Shared	**191**	**27.6x**	**Optimal for small/uniform graphs**

Table 3: BFS Performance on VRAM-Resident Graph (20M nodes)

The results lead to a critical insight: **Complexity implies Overhead**. On this smaller, uniform-degree graph (Avg=10) that fits in memory:

- **V1 (Dynamic)** outperforms V3.3 because the simple global atomic queue has minimal instruction overhead compared to the complex hybrid logic and Varint decompression of V3.3.
- **V3.1 (Shared)** is the fastest (191ms) because it combines the low overhead of V1 with Warp-Aggregation, reducing the number of global atomic adds by ≈30x compared to V1.
- **V3.3 (Compressed)** is slower here (381ms) because bandwidth is *not* the bottleneck (data is in VRAM), so the compute cost of decompression becomes a net negative.

This confirms that our advanced optimizations (Compression, Hybrid) are specifically targeted at **Massive, Sparse, Power-Law Graphs** (like Friendster) where memory bandwidth and load imbalance are the primary constraints. For smaller uniform graphs, simpler solvers (V1/V3.1) are superior.

Afforest remains bandwidth-bound due to its random atomic linking pattern, which cannot exploit the structured traversal of BFS.

5 Afforest Algorithm

We implement the Afforest algorithm for connected components:

1. **Sampling Phase** (Optional): Sample random edges to quickly link large components.
2. **Hook Phase**: Iterate over all edges, linking components via `atomicMin`.
3. **Compression**: Flatten component tree to root pointer.

For compressed graphs, we disable sampling and GCC pruning (as component IDs cannot be reliably estimated). Runtime: 21.6s—primarily bandwidth-bound by the random atomic pattern.

6 Conclusion

This project demonstrates that billion-scale graph traversal is achievable in sub-second time on consumer GPUs through careful algorithmic optimization:

- **831ms BFS on Friendster**: A 26x speedup over sequential CPU, enabled by Direct Queue Emission.
- **Warp-Level Aggregation**: Reduces atomic contention by 32x.
- **Zero-Copy + Compression**: Enables out-of-core processing with 37% bandwidth reduction.

The key insight is that **algorithmic complexity matters more than raw bandwidth**. Eliminating the O(N) distance scan in V3.3 provided a larger improvement than all previous compression optimizations combined.

Source Code: https://github.com/EvangelosMoschou/PkDSProject3.git